



Part 1: Profiling

First, we save the code as a .py file and make the following change to prevent it from being killed due to a high number of processes:

```
df = pd.concat([us_df, ca_df], ignore_index=True).sample(1000, random_state=42).reset_index(drop=True)
```

Then, we run the command below to analyze the time per function call:

```
python3 -m cProfile -s tottime ECS_CA4_P3.py > cProfile_stats.txt
```

Open  

cProfile_stats.txt
~/Desktop/ECS/CA4/P3

```
1 [EMBEDDING][INFO]: Loading model all-MiniLM-L6-v2...
2 [EMBEDDING][INFO]: Embedding column 'title'...
3 [EMBEDDING][INFO]: Embedding column 'tags'...
4 [EMBEDDING][INFO]: Embedding column 'description'...
5 Preprocessing complete. Final shape: (1000, 1181)
6 2362319229 function calls (2358581797 primitive calls) in 1044.928 seconds
7
8 Ordered by: internal time
9
10 ncalls tottime percall cumtime percall filename:lineno(function)
11 1182000 479.215 0.000 857.512 0.001 managers.py:958(fast_xs)
12 1400364159 214.585 0.000 234.423 0.000 blocks.py:1319(iget)
13 677174536 50.289 0.000 50.289 0.000 blocks.py:268(mgr_locs)
14 1182000 47.838 0.000 47.839 0.000 managers.py:984(<listcomp>)
15 3552 44.780 0.013 44.780 0.013 {built-in method torch._C._nn.linear}
16 1 16.498 16.498 1044.941 1044.941 ECS_CA4_P3.py:1(<module>)
17 1183159 10.807 0.000 32.248 0.000 cast.py:1442(find_common_type)
18 1183927 8.651 0.000 8.651 0.000 {built-in method fromkeys}
19 576 6.752 0.012 6.752 0.012 {built-in method torch._C._nn.scaled_dot_product_attention}
20 44564836/44561698 6.618 0.000 9.846 0.000 {built-in method builtins.isinstance}
21 1182007 6.319 0.000 883.650 0.001 frame.py:3988(_ixs)
22 1180000 5.877 0.000 10.255 0.000 datetimes.py:547(_box_func)
23 2321159 4.674 0.000 6.957 0.000 managers.py:1012(iget)
24 2320000 4.578 0.000 38.979 0.000 indexing.py:2529(__setitem__)
25 1163155 4.272 0.000 4.472 0.000 cast.py:1401(np_find_common_type)
26 3505332 4.030 0.000 4.473 0.000 base.py:3784(get_loc)
27 1182039 3.896 0.000 5.676 0.000 generic.py:6255(__finalize__)
28 1183152 3.740 0.000 898.087 0.001 indexing.py:1719(_getitem_axis)
29 21 3.657 0.174 3.657 0.174 {method 'read' of '_ssl._SSLSocket' objects}
30 15771253/12193957 3.467 0.000 5.311 0.000 {built-in method builtins.len}
31 2320000 3.450 0.000 30.047 0.000 frame.py:4545(_set_value)
32 2322312 3.382 0.000 4.432 0.000 cast.py:1778(np_can_hold_element)
33 1190222 3.360 0.000 3.360 0.000 {built-in method numpy.empty}
34 1180005 3.269 0.000 3.269 0.000 {method 'view' of 'numpy.generic' objects}
35 576 3.235 0.006 3.235 0.006 {built-in method torch._C._nn.gelu}
36 3547242 3.173 0.000 3.184 0.000 generic.py:6320(__setattr__)
37 1180004 3.106 0.000 3.221 0.000 utils.py:419(check_array_indexer)
38 1180004 3.027 0.000 17.422 0.000 _mixins.py:278(_getitem__)
39 1183000 2.656 0.000 12.201 0.000 series.py:1104(_getitem__)
40 1183152 2.442 0.000 902.865 0.001 indexing.py:1176(_getitem__)
41 2321152 2.404 0.000 22.459 0.000 managers.py:1298(column_setitem)
42 2320000 2.353 0.000 9.168 0.000 base.py:341(setitem_inplace)
43 1184351 2.286 0.000 3.964 0.000 blocks.py:2789(new_block)
44 1183234 2.275 0.000 2.949 0.000 generic.py:278(__init__)
45 1180004 2.224 0.000 19.839 0.000 datetimelike.py:390(__getitem__)
46 2320000 2.212 0.000 43.243 0.000 indexing.py:2577(_setitem__)
47 2402361/2399878 2.107 0.000 5.814 0.000 {built-in method builtins.any}
48 1183000 2.005 0.000 7.041 0.000 series.py:1229(_get_value)
49 4684000 1.928 0.000 2.874 0.000 managers.py:291(arrays)
50 2320000 1.924 0.000 11.833 0.000 managers.py:2021(setitem_inplace)
51 7009635 1.912 0.000 2.698 0.000 common.py:372(apply_if_callable)
52 6960000 1.825 0.000 3.271 0.000 indexing.py:2531(<genexpr>)
53 4684000 1.712 0.000 4.586 0.000 base.py:332(array)
54 1182007 1.683 0.000 8.181 0.000 frame.py:678(_constructor_sliced_from_mgr)
55 2369617 1.645 0.000 2.681 0.000 indexing.py:2765(check_dict_or_set_indexers)
56 7110465 1.548 0.000 2.561 0.000 cast.py:1472(<genexpr>)
57 4757517 1.462 0.000 1.902 0.000 generic.py:37(_check)
```

2.36 billion function calls happened in 1045 seconds.
managers.py:958(fast_xs) (used in DataFrame access, especially .iloc rows) took 479.215 seconds to run for 1182000 times and is the most time-consuming function.

Then, we run the command below to analyze the time per line:

kernprof -l -v ECS_CA4_P3.py

```
Preprocessing complete. Final shape: (1000, 1181)
Wrote profile results to ECS_CA4_P3.py.lprof
Timer unit: 1e-06 s

Total time: 1866.92 s
File: ECS_CA4_P3.py
Function: run_pipeline at line 8
```

Line #	Hits	Time	Per Hit	% Time	Line Contents
8					@profile
9					def run_pipeline():
10	1	2.0	2.0	0.0	TEXTUAL_COLUMNS = ["title", "tags", "description"]
11	1	1.0	1.0	0.0	EMBEDDING_MODEL = "all-MiniLM-L6-v2"
12	1	1.0	1.0	0.0	EMBEDDING_DIM = 384
13	1	1.0	1.0	0.0	OUTPUT_DIR = "tmp/embeddings/"
14	1	26.0	26.0	0.0	os.makedirs(OUTPUT_DIR, exist_ok=True)
15					
16	1	396619.0	396619.0	0.0	us_df = pd.read_csv("USvideos.csv")
17	1	1264.0	1264.0	0.0	us_df["country"] = "US"
18					
19	1	486403.0	486403.0	0.0	ca_df = pd.read_csv("CAvideos.csv")
20	1	604.0	604.0	0.0	ca_df["country"] = "CA"
21					
22	1	8895.0	8895.0	0.0	df = pd.concat([us_df, ca_df], ignore_index=True).sample(1000, random_state=42).reset_index(drop=True)
23					
24	1	64.0	64.0	0.0	print(f"[EMBEDDING][INFO]: Loading model {EMBEDDING_MODEL}...")
25	1	4592619.0	4592619.0	0.2	model = SentenceTransformer(EMBEDDING_MODEL)
26					
27	1	4.0	4.0	0.0	def clean_tags(text):
28					return " ".join(tag.replace("'", '') for tag in str(text).split(' '))
29					
30	4	11.0	2.8	0.0	for col in TEXTUAL_COLUMNS:
31	3	247.0	82.3	0.0	print(f"[EMBEDDING][INFO]: Embedding column '{col}'...")
32	3	6.0	2.0	0.0	if col == "tags":
33	1	8191.0	8191.0	0.0	text_data = df[col].fillna("").apply(clean_tags).tolist()
34					else:
35	2	2449.0	1224.5	0.0	text_data = df[col].fillna("").astype(str).tolist()
36					
37	3	54372205.0	18124068.3	2.9	emb = model.encode(text_data, show_progress_bar=True, batch_size=32)
38	3	1770.0	590.0	0.0	emb_df = pd.DataFrame(emb, columns=[f"{col}_emb_{i}" for i in range(emb.shape[1])])
39	3	6551.0	2183.7	0.0	df = pd.concat([df.reset_index(drop=True), emb_df], axis=1)
40					
41	1	2.0	2.0	0.0	def count_tags_loop(tag_str):
42					if pd.isna(tag_str):
43					return 0
44					count = 0
45					for tag in tag_str.split(" "):
46					if tag.strip() != "":
47					count += 1
48					return count
49					
50	1	1.0	1.0	0.0	tag_counts = []

```

51      1001      1453.0      1.5      0.0      for i in range(len(df)):
52      1000      898718.0      898.7      0.0          tag_counts.append(count_tags_loop(df.iloc[i]["tags"]))
53      1      893.0      893.0      0.0          df["tag_count"] = tag_counts
54
55      1      2.0      2.0      0.0          publish_dates = []
56      1      1.0      1.0      0.0          publish_hours = []
57      1001      1101.0      1.1      0.0          for i in range(len(df)):
58      1000      1032.0      1.0      0.0              try:
59      1000      922827.0      922.8      0.0                  dt = datetime.strptime(df.iloc[i]["publish_time"], "%Y-%m-%dT%H:%M:%S.%fZ")
60      1000      1482.0      1.5      0.0                  publish_dates.append(dt)
61      1000      1331.0      1.3      0.0                  publish_hours.append(dt.hour)
62                  except Exception:
63                      publish_dates.append(pd.NaT)
64                      publish_hours.append(np.nan)
65
66      1      3991.0      3991.0      0.0          df["publish_time"] = publish_dates
67      1      571.0      571.0      0.0          df["publish_hour"] = publish_hours
68
69      4      5.0      1.2      0.0          for col in TEXTUAL_COLUMNS:
70      3      157.0      52.3      0.0              if col in df.columns:
71      3      1153.0      384.3      0.0                  del df[col]
72
73      1      1.0      1.0      0.0          engagement_rates = []
74      1      1.0      1.0      0.0          ratios = []
75      1001      1126.0      1.1      0.0          for i in range(len(df)):
76      1000      893254.0      893.3      0.0              row = df.iloc[i]
77      1000      10914.0      10.9      0.0              views = row["views"]
78      1000      6724.0      6.7      0.0              likes = row["likes"]
79      1000      6212.0      6.2      0.0              dislikes = row["dislikes"]
80      1000      6211.0      6.2      0.0              comments = row["comment_count"]
81      1000      2003.0      2.0      0.0              engagement_rates.append((likes + dislikes + comments) / (views + 1))
82      1000      1449.0      1.4      0.0              ratios.append(likes / (dislikes + 1))
83      1      754.0      754.0      0.0          df["engagement_rate"] = engagement_rates
84      1      678.0      678.0      0.0          df["like_dislike_ratio"] = ratios
85
86      1      1611.0      1611.0      0.0          unique_cats = sorted(df["category_id"].dropna().unique())
87      1      2.0      2.0      0.0          one_hot = []
88      1001      1207.0      1.2      0.0          for i in range(len(df)):
89      1000      1117.0      1.1      0.0              row = []
90      18000      27150.0      1.5      0.0              for cat in unique_cats:
91      17000      15600288.0      917.7      0.8                  row.append(1 if df.iloc[i]["category_id"] == cat else 0)
92      1000      1406.0      1.4      0.0              one_hot.append(row)
93
94      1      2863.0      2863.0      0.0          cat_df = pd.DataFrame(one_hot, columns=[f"cat_{int(c)}" for c in unique_cats])
95      1      3698.0      3698.0      0.0          df = pd.concat([df.reset_index(drop=True), cat_df], axis=1)
96      1      1473.0      1473.0      0.0          df = df.drop(columns=["category_id"])
97
98      1      2.0      2.0      0.0          bool_cols = ["comments_disabled", "ratings_disabled", "video_error_or_removed"]
99      4      8.0      2.0      0.0          for col in bool_cols:
100      3      2061.0      687.0      0.0              df[col] = [int(val) for val in df[col]]
101      1      1352.0      1352.0      0.0          df = df.drop(columns=bool_cols)
102
103      1      2.0      2.0      0.0          seen_rows = set()
104      1      1.0      1.0      0.0          deduped_rows = []

```

```

105      1001      1562.0      1.6      0.0          for i in range(len(df)):
106      1000      981842.0      981.8      0.1              row_tuple = tuple(df.iloc[i].values)
107      1000      19336.0      19.3      0.0              if row_tuple not in seen_rows:
108      1000      17629.0      17.6      0.0                  seen_rows.add(row_tuple)
109      1000      953327.0      953.3      0.1                  deduped_rows.append(df.iloc[i])
110      1      172910.0      172910.0      0.0          df = pd.DataFrame(deduped_rows).reset_index(drop=True)
111
112      1      3.0      3.0      0.0          numeric_attributes = [
113      "views", "publish_hour", "likes", "dislikes", "comment_count",
114      "engagement_rate", "like_dislike_ratio", "tag_count"
115      ]
116      1      235.0      235.0      0.0          numeric_attributes += [col for col in df.columns if "_emb_" in col]
117
118      1161      2563.0      2.2      0.0          for col in numeric_attributes:
119      1160      8043.0      6.9      0.0              transformed = []
120      1161160      1841835.0      1.6      0.1              for i in range(len(df)):
121      1160000      1730356167.0      1491.7      92.7                  transformed.append(np.log1p(df.iloc[i][col]))
122      1160      659873.0      568.9      0.0              df[col] = transformed
123
124      1      49.0      49.0      0.0          minmax_scaler = MinMaxScaler()
125      1      76337.0      76337.0      0.0          scaled_minmax = minmax_scaler.fit_transform(df[numeric_attributes])
126      1161      1980.0      1.7      0.0          for j, col in enumerate(numeric_attributes):
127      1161160      1270199.0      1.1      0.1              for i in range(len(df)):
128      1160000      25483031.0      22.0      1.4                  df.at[i, col] = np.float32(scaled_minmax[i][j])
129
130      1      10.0      10.0      0.0          standard_scaler = StandardScaler()
131      1      72600.0      72600.0      0.0          scaled_standard = standard_scaler.fit_transform(df[numeric_attributes])
132      1161      1930.0      1.7      0.0          for j, col in enumerate(numeric_attributes):
133      1161160      1268012.0      1.1      0.1              for i in range(len(df)):
134      1160000      25421565.0      21.9      1.4                  df.at[i, col] = np.float32(scaled_standard[i][j])
135
136      1      23007.0      23007.0      0.0          df = df.drop(columns=["likes", "dislikes"])
137      1      103.0      103.0      0.0          print("Preprocessing complete. Final shape:", df.shape)

```

Line 121 (transformed.append(np.log1p(df.iloc[i][col]))), which took 1730.36 seconds to run, is the most time-consuming one. This line accounts for 92.7% of the total execution time.

Then, we run the command below to analyze the memory usage per line:

kernprof -l -v ECS_CA4_P3.py

Line #	Mem usage	Increment	Line Contents
=====			
8	651.629 MiB	651.629 MiB	@profile
9			def run_pipeline():
10	651.629 MiB	0.000 MiB	TEXTUAL_COLUMNS = ["title", "tags", "description"]
11	651.629 MiB	0.000 MiB	EMBEDDING_MODEL = "all-MiniLM-L6-v2"
12	651.629 MiB	0.000 MiB	EMBEDDING_DIM = 384
13	651.629 MiB	0.000 MiB	OUTPUT_DIR = "tmp/embeddings/"
14	651.629 MiB	0.000 MiB	os.makedirs(OUTPUT_DIR, exist_ok=True)
15			
16	687.438 MiB	35.809 MiB	us_df = pd.read_csv("USvideos.csv")
17	687.562 MiB	0.125 MiB	us_df["country"] = "US"
18			
19	730.297 MiB	42.734 MiB	ca_df = pd.read_csv("CAvideos.csv")
20	730.297 MiB	0.000 MiB	ca_df["country"] = "CA"
21			
22	731.797 MiB	1.500 MiB	df = pd.concat([us_df, ca_df], ignore_index=True).sample(1000, random_state=42).reset_index(drop=True)
23			
24	731.797 MiB	0.000 MiB	print(f"[EMBEDDING][INFO]: Loading model {EMBEDDING_MODEL}...")
25	756.902 MiB	25.105 MiB	model = SentenceTransformer(EMBEDDING_MODEL)
26			
27	945.816 MiB	0.000 MiB	def clean_tags(text):
28	945.816 MiB	0.000 MiB	return " ".join(tag.replace("'", '') for tag in str(text).split(' '))
29			
30	1001.566 MiB	0.000 MiB	for col in TEXTUAL_COLUMNS:
31	982.613 MiB	0.000 MiB	print(f"[EMBEDDING][INFO]: Embedding column '{col}'...")
32	982.613 MiB	0.000 MiB	if col == "tags":
33	945.816 MiB	0.000 MiB	text_data = df[col].fillna("").apply(clean_tags).tolist()
34			else:
35	982.738 MiB	0.125 MiB	text_data = df[col].fillna("").astype(str).tolist()
36			
37	1001.566 MiB	244.539 MiB	emb = model.encode(text_data, show_progress_bar=True, batch_size=32)
38	1001.566 MiB	0.000 MiB	emb_df = pd.DataFrame(emb, columns=[f"{col}_emb_{i}" for i in range(emb.shape[1])])
39	1001.566 MiB	0.000 MiB	df = pd.concat([df.reset_index(drop=True), emb_df], axis=1)
40			
41	1001.566 MiB	0.000 MiB	def count_tags_loop(tag_str):
42	1001.566 MiB	0.000 MiB	if pd.isna(tag_str):
43			return 0
44	1001.566 MiB	0.000 MiB	count = 0
45	1001.566 MiB	0.000 MiB	for tag in tag_str.split(" "):
46	1001.566 MiB	0.000 MiB	if tag.strip() != "":
47	1001.566 MiB	0.000 MiB	count += 1
48	1001.566 MiB	0.000 MiB	return count
49			
50	1001.566 MiB	0.000 MiB	tag_counts = []
51	1001.566 MiB	0.000 MiB	for i in range(len(df)):
52	1001.566 MiB	0.000 MiB	tag_counts.append(count_tags_loop(df.iloc[i]["tags"]))
53	1001.566 MiB	0.000 MiB	df["tag_count"] = tag_counts
54			
55	1001.566 MiB	0.000 MiB	publish_dates = []
56	1001.566 MiB	0.000 MiB	publish_hours = []
57	1001.566 MiB	0.000 MiB	for i in range(len(df)):
58	1001.566 MiB	0.000 MiB	try:
59	1001.566 MiB	0.000 MiB	dt = datetime.strptime(df.iloc[i]["publish_time"], "%Y-%m-%dT%H:%M:%S.%fZ")
60	1001.566 MiB	0.000 MiB	publish_dates.append(dt)
61	1001.566 MiB	0.000 MiB	publish_hours.append(dt.hour)

```

62         except Exception:
63             publish_dates.append(pd.NaT)
64             publish_hours.append(np.nan)
65
66 1003.066 MiB    1.500 MiB    df["publish_time"] = publish_dates
67 1003.066 MiB    0.000 MiB    df["publish_hour"] = publish_hours
68
69 1003.191 MiB    0.000 MiB    for col in TEXTUAL_COLUMNS:
70 1003.191 MiB    0.000 MiB        if col in df.columns:
71 1003.191 MiB    0.125 MiB            del df[col]
72
73 1003.191 MiB    0.000 MiB    engagement_rates = []
74 1003.191 MiB    0.000 MiB    ratios = []
75 1003.316 MiB    0.000 MiB    for i in range(len(df)):
76 1003.316 MiB    0.125 MiB        row = df.iloc[i]
77 1003.316 MiB    0.000 MiB        views = row["views"]
78 1003.316 MiB    0.000 MiB        likes = row["likes"]
79 1003.316 MiB    0.000 MiB        dislikes = row["dislikes"]
80 1003.316 MiB    0.000 MiB        comments = row["comment_count"]
81 1003.316 MiB    0.000 MiB        engagement_rates.append((likes + dislikes + comments) / (views + 1))
82 1003.316 MiB    0.000 MiB        ratios.append(likes / (dislikes + 1))
83 1003.316 MiB    0.000 MiB    df["engagement_rate"] = engagement_rates
84 1003.316 MiB    0.000 MiB    df["like_dislike_ratio"] = ratios
85
86 1003.691 MiB    0.375 MiB    unique_cats = sorted(df["category_id"].dropna().unique())
87 1003.691 MiB    0.000 MiB    one_hot = []
88 1003.691 MiB    0.000 MiB    for i in range(len(df)):
89 1003.691 MiB    0.000 MiB        row = []
90 1003.691 MiB    0.000 MiB        for cat in unique_cats:
91 1003.691 MiB    0.000 MiB            row.append(1 if df.iloc[i]["category_id"] == cat else 0)
92 1003.691 MiB    0.000 MiB        one_hot.append(row)
93
94 1003.691 MiB    0.000 MiB    cat_df = pd.DataFrame(one_hot, columns=[f"cat_{int(c)}" for c in unique_cats])
95 1003.691 MiB    0.000 MiB    df = pd.concat([df.reset_index(drop=True), cat_df], axis=1)
96 1003.816 MiB    0.125 MiB    df = df.drop(columns=["category_id"])
97
98 1003.816 MiB    0.000 MiB    bool_cols = ["comments_disabled", "ratings_disabled", "video_error_or_removed"]
99 1003.816 MiB    0.000 MiB    for col in bool_cols:
100 1003.816 MiB    0.000 MiB        df[col] = [int(val) for val in df[col]]
101 1003.816 MiB    0.000 MiB    df = df.drop(columns=bool_cols)
102
103 1003.816 MiB    0.000 MiB    seen_rows = set()
104 1003.816 MiB    0.000 MiB    deduped_rows = []
105 1076.566 MiB    0.000 MiB    for i in range(len(df)):
106 1076.566 MiB    36.375 MiB        row_tuple = tuple(df.iloc[i].values)
107 1076.566 MiB    0.125 MiB        if row_tuple not in seen_rows:
108 1076.566 MiB    0.000 MiB            seen_rows.add(row_tuple)
109 1076.566 MiB    36.250 MiB            deduped_rows.append(df.iloc[i])
110 1077.316 MiB    0.750 MiB    df = pd.DataFrame(deduped_rows).reset_index(drop=True)
111
112 1077.316 MiB    0.000 MiB    numeric_attributes = [
113        "views", "publish_hour", "likes", "dislikes", "comment_count",
114        "engagement_rate", "like_dislike_ratio", "tag_count"
115    ]
116 1077.316 MiB    0.000 MiB    numeric_attributes += [col for col in df.columns if "_emb_" in col]
117
118 1078.566 MiB    0.000 MiB    for col in numeric_attributes:
119 1078.566 MiB    0.000 MiB        transformed = []
120 1078.566 MiB    0.000 MiB        for i in range(len(df)):
121 1078.566 MiB    1.125 MiB            transformed.append(np.log1p(df.iloc[i][col]))
122 1078.566 MiB    0.125 MiB        df[col] = transformed
123
124 1078.566 MiB    0.000 MiB    minmax_scaler = MinMaxScaler()
125 1080.191 MiB    1.625 MiB    scaled_minmax = minmax_scaler.fit_transform(df[numeric_attributes])
126 1080.316 MiB    0.000 MiB    for j, col in enumerate(numeric_attributes):
127 1080.316 MiB    0.000 MiB        for i in range(len(df)):
128 1080.316 MiB    0.125 MiB            df.at[i, col] = np.float32(scaled_minmax[i][j])
129
130 1080.316 MiB    0.000 MiB    standard_scaler = StandardScaler()
131 1080.316 MiB    0.000 MiB    scaled_standard = standard_scaler.fit_transform(df[numeric_attributes])
132 1080.316 MiB    0.000 MiB    for j, col in enumerate(numeric_attributes):
133 1080.316 MiB    0.000 MiB        for i in range(len(df)):
134 1080.316 MiB    0.000 MiB            df.at[i, col] = np.float32(scaled_standard[i][j])
135
136 1080.316 MiB    0.000 MiB    df = df.drop(columns=["likes", "dislikes"])
137 1080.316 MiB    0.000 MiB    print("Preprocessing complete. Final shape:", df.shape)

```


Initial memory usage is 651 MiB and Peak memory usage is 1080 MiB.

Then, we run the command below to analyze the real-time CPU usage:

py-spy top -- python3 ECS_CA4_P3.py

```
Collecting samples from 'python3 ECS_CA4_P3.py' (python v3.10.12)
Total Samples 6100
GIL: 1.00%, Active: 99.00%, Threads: 2
```

%Own	%Total	OwnTime	TotalTime	Function (filename)
68.00%	68.00%	38.71s	38.71s	forward (torch/nn/modules/linear.py)
22.00%	97.00%	6.40s	49.69s	forward (transformers/models/bert/modeling_bert.py)
3.00%	3.00%	3.00s	3.00s	forward (transformers/activations.py)
0.00%	0.00%	0.860s	0.950s	read (pandas/io/parsers/c_parser_wrapper.py)
0.00%	0.00%	0.850s	4.54s	_call_with_frames_removed (<frozen importlib._bootstrap>)
0.00%	0.00%	0.650s	0.650s	get_data (<frozen importlib._bootstrap_external>)
4.00%	4.00%	0.630s	0.630s	layer_norm (torch/nn/functional.py)
0.00%	0.00%	0.550s	0.550s	_compile_bytecode (<frozen importlib._bootstrap_external>)
0.00%	0.00%	0.490s	0.820s	create_import_structure_from_path (transformers/utils/import_utils.py)
0.00%	46.00%	0.380s	30.24s	apply_chunking_to_forward (transformers/pytorch_utils.py)
0.00%	0.00%	0.270s	0.270s	fetch_all__ (transformers/utils/import_utils.py)
0.00%	0.00%	0.260s	0.260s	embedding (torch/nn/functional.py)
0.00%	0.00%	0.230s	0.230s	read (ssl.py)
0.00%	0.00%	0.170s	0.180s	_path_hooks (<frozen importlib._bootstrap_external>)
1.00%	1.00%	0.160s	0.280s	_batch_encode_plus (transformers/tokenization_utils_fast.py)
0.00%	0.00%	0.130s	0.130s	_create_fn (dataclasses.py)
1.00%	1.00%	0.120s	0.120s	forward (sentence_transformers/models/Pooling.py)
0.00%	0.00%	0.110s	0.110s	_path_stat (<frozen importlib._bootstrap_external>)
0.00%	98.00%	0.090s	49.81s	_call_impl (torch/nn/modules/module.py)
0.00%	0.00%	0.080s	0.080s	__setattr__ (enum.py)
0.00%	0.00%	0.070s	0.070s	impl (torch/library.py)
0.00%	0.00%	0.070s	0.070s	as_tensor (transformers/tokenization_utils_base.py)
0.00%	0.00%	0.060s	0.070s	_joinrealpath (posixpath.py)
0.00%	0.00%	0.060s	0.060s	decode (codecs.py)
0.00%	0.00%	0.060s	0.060s	__init__ (ctypes/_init_.py)
0.00%	0.00%	0.050s	0.060s	dedent (textwrap.py)
0.00%	0.00%	0.050s	0.050s	<listcomp> (sentence_transformers/SentenceTransformer.py)
0.00%	0.00%	0.050s	0.060s	__init__ (inspect.py)
0.00%	0.00%	0.050s	0.060s	ssl_wrap_socket (urllib3/util/ssl_.py)
0.00%	0.00%	0.050s	0.060s	_has_script_object_arg (torch/_ops.py)
0.00%	0.00%	0.040s	0.040s	_convert_encoding (transformers/tokenization_utils_fast.py)
0.00%	0.00%	0.040s	0.040s	transpose_for_scores (transformers/models/bert/modeling_bert.py)
0.00%	0.00%	0.040s	0.050s	_merge_blocks (pandas/core/internals/managers.py)
0.00%	0.00%	0.030s	0.030s	join (posixpath.py)
0.00%	0.00%	0.030s	0.110s	_signature_from_callable (inspect.py)
0.00%	0.00%	0.030s	0.030s	__getattr__ (torch/nn/modules/module.py)

In the example screenshot, total samples are 6100. GIL (Global Interpreter Lock) is 1.00% that means 99% of time is spent in native code not blocked by Python itself. Active is 99.00% that means the program is CPU-active almost all the time. Threads is 2 that means some operations use background threads.

Part 2: Optimization

1. `.iloc[i]` is extremely slow and memory-intensive.

```
tag_counts = []
for i in range(len(df)):
    tag_counts.append(count_tags_loop(df.iloc[i]["tags"]))
df["tag_count"] = tag_counts
```



```
df["tag_count"] = df["tags"].fillna("").apply(lambda x: len([t for t in
str(x).split("|") if t.strip()])))
```

2. `.iloc[i]` is extremely slow and memory-intensive.

```
for i in range(len(df)):
    row = df.iloc[i]
    views = row["views"]
    likes = row["likes"]
    dislikes = row["dislikes"]
    comments = row["comment_count"]
    engagement_rates.append((likes + dislikes + comments) / (views + 1))
    ratios.append(likes / (dislikes + 1))
```



```
df["engagement_rate"] = (df["likes"] + df["dislikes"] + df["comment_count"]) /
(df["views"] + 1)
df["like_dislike_ratio"] = df["likes"] / (df["dislikes"] + 1)
```

3. The code manually loops over rows and categories, creating one-hot encodings with conditionals.

```
unique_cats = sorted(df["category_id"].dropna().unique())
one_hot = []
for i in range(len(df)):
    row = []
    for cat in unique_cats:
        row.append(1 if df.iloc[i]["category_id"] == cat else 0)
    one_hot.append(row)
cat_df = pd.DataFrame(one_hot, columns=[f"cat_{int(c)}" for c in
unique_cats])
df = pd.concat([df.reset_index(drop=True), cat_df], axis=1)
```



```
cat_df = pd.get_dummies(df["category_id"], prefix="cat")
df = pd.concat([df, cat_df], axis=1)
```


cProfile:

```
1 [EMBEDDING][INFO]: Loading model all-MiniLM-L6-v2...
2 [EMBEDDING][INFO]: Embedding column 'title'...
3 [EMBEDDING][INFO]: Embedding column 'tags'...
4 [EMBEDDING][INFO]: Embedding column 'description'...
5 Preprocessing complete. Final shape: (1000, 1181)
6 11148287 function calls (10944247 primitive calls) in 68.517 seconds
7
8 Ordered by: internal time
9
10 ncalls tottime percall cumtime percall filename:lineno(function)
11 3552 43.161 0.012 43.161 0.012 {built-in method torch._C.nn.linear}
12 576 6.508 0.011 6.508 0.011 {built-in method torch._C.nn.scaled_dot_product_attention}
13 576 3.459 0.006 3.459 0.006 {built-in method torch._C.nn.gelu}
14 21 3.438 0.164 3.438 0.164 {method 'read' of '_ssl.SSLSocket' objects}
15 1248 0.783 0.001 0.783 0.001 {built-in method torch.layer_norm}
16 2 0.766 0.383 0.836 0.418 c_parser_wrapper.py:222(read)
17 96 0.570 0.006 0.570 0.006 {method 'encode_batch' of 'tokenizers.Tokenizer' objects}
18 576 0.563 0.001 34.167 0.059 pytorch_utils.py:175(apply_chunking_to_forward)
19 576 0.409 0.001 18.715 0.032 modeling_bert.py:506(forward)
20 192/191 0.379 0.002 0.381 0.002 {built-in method imp.create_dynamic}
21 576 0.343 0.001 14.799 0.026 modeling_bert.py:519(forward)
22 576 0.334 0.001 4.377 0.008 modeling_bert.py:433(forward)
23 3696 0.284 0.000 0.590 0.000 import_utils.py:2309(fetch_all_)
24 4097 0.262 0.000 0.262 0.000 {built-in method marshal.loads}
25 2940820 0.215 0.000 0.215 0.000 {method 'startswith' of 'str' objects}
26 288 0.187 0.001 0.187 0.001 {built-in method torch.embedding}
27 7103/1 0.182 0.000 68.529 68.529 {built-in method builtins.exec}
28 5884 0.174 0.000 0.181 0.000 generic.py:6255(_finalize_)
29 3937 0.148 0.000 0.170 0.000 {method 'read' of '_io.TextIOWrapper' objects}
30 1064/319 0.142 0.000 1.034 0.003 import_utils.py:2351(create_import_structure_from_path)
31 11136/288 0.133 0.000 56.686 0.197 module.py:1755(_call_impl)
32 13920 0.131 0.000 0.131 0.000 {method 'splitlines' of 'str' objects}
33 2973 0.118 0.000 0.189 0.000 SentenceTransformer.py:1999(<listcomp>)
34 36447 0.116 0.000 0.117 0.000 {built-in method posix.stat}
35 1674969/1641076 0.108 0.000 0.116 0.000 {built-in method builtins.len}
36 15253 0.099 0.000 0.103 0.000 <frozen importlib._bootstrap:100(acquire)
37 1 0.099 0.099 0.099 0.099 {built-in method _socket.getaddrinfo}
38 1 0.099 0.099 0.099 0.099 {method 'connect' of '_socket.socket' objects}
39 15253 0.097 0.000 0.112 0.000 <frozen importlib._bootstrap:179(_get_module_lock)
40 96 0.092 0.001 0.347 0.004 modeling_bert.py:149(forward)
41 4114 0.090 0.000 0.090 0.000 {method 'read' of '_io.BufferedReader' objects}
42 371 0.088 0.000 0.088 0.000 {built-in method torch.tensor}
43 1 0.088 0.088 0.088 0.088 {method 'do_handshake' of '_ssl.SSLSocket' objects}
44 7854/7742 0.086 0.000 0.380 0.000 {built-in method builtins._build_class_}
45 834758/824491 0.085 0.000 0.112 0.000 {built-in method builtins.isinstance}
46 5037 0.082 0.000 0.082 0.000 {built-in method _codecs.utf_8.decode}
47 1382 0.068 0.000 0.068 0.000 {method 'to' of 'torch._C.TensorBase' objects}
48 6633 0.065 0.000 0.077 0.000 functools.py:35(update_wrapper)
49 4099 0.064 0.000 0.064 0.000 {built-in method io.open_code}
50 96 0.053 0.001 0.137 0.001 Pooling.py:135(forward)
51 7293/666 0.052 0.000 4.973 0.007 <frozen importlib._bootstrap:1022(_find_and_load)
52 3 0.050 0.017 57.747 19.249 SentenceTransformer.py:826(encode)
53 217017/216897 0.050 0.000 0.153 0.000 {built-in method builtins.getattr}
54 1 0.048 0.048 0.048 0.048 {method 'load_verify_locations' of '_ssl.SSLContext' objects}
55 96 0.046 0.000 56.050 0.584 modeling_bert.py:619(forward)
```

Before optimization, *managers.py:958(fast_xs)* and *blocks.py:1319(iget)* dominated CPU time due to inefficient DataFrame operations (e.g., *iloc*, *.at[]*, manual loops).

Log transform, MinMaxScaler, and StandardScaler were applied with for loops and *.at[i, j]* style updates which were extremely slow and memory-inefficient.

After optimization, these high-cost pandas internal calls are gone or negligible because of using vectorized pandas operations, avoiding slow row-by-row *.iloc* loops.

Transformed with *apply* + vectorized operations, and scalers applied directly to the entire DataFrame slice.

line-profiler:

```
Preprocessing complete. Final shape: (1000, 1181)
Wrote profile results to ECS_CA4_P3_optimized.py.lprof
Timer unit: 1e-06 s

Total time: 68.027 s
File: ECS_CA4_P3_optimized.py
Function: run_pipeline at line 8

Line #      Hits          Time  Per Hit    % Time  Line Contents
=====
8           1           3.0      3.0      0.0      @profile
9           1           1.0      1.0      0.0      def run_pipeline():
10          1           1.0      1.0      0.0          TEXTUAL_COLUMNS = ["title", "tags", "description"]
11          1           1.0      1.0      0.0          EMBEDDING_MODEL = "all-MiniLM-L6-v2"
12          1           1.0      1.0      0.0          EMBEDDING_DIM = 384
13          1           1.0      1.0      0.0          OUTPUT_DIR = "tmp/embeddings/"
14          1          28.0     28.0      0.0          os.makedirs(OUTPUT_DIR, exist_ok=True)
15
16          1      438056.0  438056.0      0.6          us_df = pd.read_csv("USvideos.csv")
17          1      1223.0    1223.0      0.0          us_df["country"] = "US"
18          1      576184.0  576184.0      0.8          ca_df = pd.read_csv("CAvideos.csv")
19          1      545.0     545.0      0.0          ca_df["country"] = "CA"
20
21          1       7002.0   7002.0      0.0          df = pd.concat([us_df, ca_df], ignore_index=True).sample(1000, random_state=42).reset_index(drop=True)
22
23          1       123.0     123.0      0.0          print(f"[EMBEDDING][INFO]: Loading model {EMBEDDING_MODEL}...")
24          1      4505658.0  4505658.0      6.6          model = SentenceTransformer(EMBEDDING_MODEL)
25
26          1           3.0      3.0      0.0          def clean_tags(text):
27              1           1.0      1.0      0.0              return " ".join(tag.replace("'", '') for tag in str(text).split('|'))
28
29          4           8.0      2.0      0.0          for col in TEXTUAL_COLUMNS:
30              3          219.0     73.0      0.0              print(f"[EMBEDDING][INFO]: Embedding column '{col}'...")
31              3           4.0      1.3      0.0              if col == "tags":
32                  1          8997.0   8997.0      0.0                  text_data = df[col].fillna("").apply(clean_tags).tolist()
33              2          3465.0   1732.5      0.0              else:
34                  2          3465.0   1732.5      0.0                  text_data = df[col].fillna("").astype(str).tolist()
35
36          3      61489954.0  20496651.3      90.4              emb = model.encode(text_data, show_progress_bar=True, batch_size=32)
37          3       1945.0     648.3      0.0              emb_df = pd.DataFrame(emb, columns=[f"{col}_emb_{i}" for i in range(emb.shape[1])])
38          3       5943.0   1981.0      0.0              df = pd.concat([df.reset_index(drop=True), emb_df], axis=1)
39
40          1       5425.0   5425.0      0.0          df["tag_count"] = df["tags"].fillna("").apply(lambda x: len([t for t in str(x).split("|") if t.strip()]))
41
42          1      19194.0  19194.0      0.0          df["publish_time"] = pd.to_datetime(df["publish_time"], errors="coerce", utc=True)
43          1      1431.0   1431.0      0.0          df["publish_hour"] = df["publish_time"].dt.hour
44
45          1       4958.0   4958.0      0.0          df.drop(columns=[col for col in TEXTUAL_COLUMNS if col in df.columns], inplace=True)
46
47          1      1894.0   1894.0      0.0          df["engagement_rate"] = (df["likes"] + df["dislikes"] + df["comment_count"]) / (df["views"] + 1)
48          1       581.0     581.0      0.0          df["like_dislike_ratio"] = df["likes"] / (df["dislikes"] + 1)
49
50          1      1860.0   1860.0      0.0          cat_df = pd.get_dummies(df["category_id"], prefix="cat")
51          1      2130.0   2130.0      0.0          df = pd.concat([df, cat_df], axis=1)
52          1      6787.0   6787.0      0.0          df.drop(columns=["category_id"], inplace=True)
53
54          1           2.0      2.0      0.0          bool_cols = ["comments_disabled", "ratings_disabled", "video_error_or_removed"]
55          1      1329.0   1329.0      0.0          df[bool_cols] = df[bool_cols].astype(int)
56          1      1004.0   1004.0      0.0          df.drop(columns=bool_cols, inplace=True)
57
58          1      140663.0  140663.0      0.2          df = df.drop_duplicates().reset_index(drop=True)
59
60          2           7.0      3.5      0.0          numeric_attributes = [
61              1           1.0      1.0      0.0              "views", "publish_hour", "likes", "dislikes", "comment_count",
62              1           1.0      1.0      0.0              "engagement_rate", "like_dislike_ratio", "tag_count"
63          ] + [col for col in df.columns if "_emb_" in col]
64
65          1      486082.0  486082.0      0.7          df[numeric_attributes] = df[numeric_attributes].apply(lambda col: np.log1p(col))
66
67          1           8.0      8.0      0.0          minmax_scaler = MinMaxScaler()
68          1      138957.0  138957.0      0.2          df[numeric_attributes] = minmax_scaler.fit_transform(df[numeric_attributes])
69
70          1          10.0     10.0      0.0          standard_scaler = StandardScaler()
71          1      147623.0  147623.0      0.2          df[numeric_attributes] = standard_scaler.fit_transform(df[numeric_attributes])
72
73          1      27450.0   27450.0      0.0          df.drop(columns=["likes", "dislikes"], inplace=True)
74
75          1          66.0     66.0      0.0          print("Preprocessing complete. Final shape:", df.shape)
```

For example, from nested loops to *apply(np.log1p)* has a 30x speedup.

memory-profiler:

```
Preprocessing complete. Final shape: (1000, 1181)
Filename: ECS_CA4_P3_optimized.py

Line #      Mem usage      Increment      Line Contents
=====
8      649.922 MiB      649.922 MiB      @profile
9
10     649.922 MiB      0.000 MiB      def run_pipeline():
11     649.922 MiB      0.000 MiB          TEXTUAL_COLUMNS = ["title", "tags", "description"]
12     649.922 MiB      0.000 MiB          EMBEDDING_MODEL = "all-MiniLM-L6-v2"
13     649.922 MiB      0.000 MiB          EMBEDDING_DIM = 384
14     649.922 MiB      0.000 MiB          OUTPUT_DIR = "tmp/embeddings/"
15     649.922 MiB      0.000 MiB          os.makedirs(OUTPUT_DIR, exist_ok=True)
16
17     685.488 MiB      35.566 MiB          us_df = pd.read_csv("USvideos.csv")
18     685.738 MiB      0.250 MiB          us_df["country"] = "US"
19     753.105 MiB      67.367 MiB          ca_df = pd.read_csv("CAvideos.csv")
20     753.105 MiB      0.000 MiB          ca_df["country"] = "CA"
21
22     753.730 MiB      0.625 MiB          df = pd.concat([us_df, ca_df], ignore_index=True).sample(1000, random_state=42).reset_index(drop=True)
23
24     770.184 MiB      16.453 MiB          print(f"[EMBEDDING][INFO]: Loading model {EMBEDDING_MODEL}...")
25     770.184 MiB      0.000 MiB          model = SentenceTransformer(EMBEDDING_MODEL)
26
27     914.297 MiB      0.000 MiB          def clean_tags(text):
28     914.297 MiB      0.000 MiB              return " ".join(tag.replace("'", '') for tag in str(text).split('|'))
29
30     980.711 MiB      -24.027 MiB          for col in TEXTUAL_COLUMNS:
31     980.711 MiB      0.000 MiB              print(f"[EMBEDDING][INFO]: Embedding column '{col}'...")
32     980.711 MiB      0.000 MiB              if col == "tags":
33     914.297 MiB      0.000 MiB                  text_data = df[col].fillna("").apply(clean_tags).tolist()
34     980.836 MiB      0.125 MiB              else:
35     980.836 MiB      0.000 MiB                  text_data = df[col].fillna("").astype(str).tolist()
36
37     980.711 MiB      186.250 MiB          emb = model.encode(text_data, show_progress_bar=True, batch_size=32)
38     980.711 MiB      -9298.582 MiB          emb_df = pd.DataFrame(emb, columns=[f"{col}_emb_{i}" for i in range(emb.shape[1])])
39     980.711 MiB      -23.902 MiB          df = pd.concat([df.reset_index(drop=True), emb_df], axis=1)
40
41     956.684 MiB      -24.027 MiB          df["tag_count"] = df["tags"].fillna("").apply(lambda x: len([t for t in str(x).split("|") if t.strip()]))
42
43     958.934 MiB      2.250 MiB          df["publish_time"] = pd.to_datetime(df["publish_time"], errors="coerce", utc=True)
44     959.059 MiB      0.125 MiB          df["publish_hour"] = df["publish_time"].dt.hour
45
46     959.434 MiB      0.375 MiB          df.drop(columns=[col for col in TEXTUAL_COLUMNS if col in df.columns], inplace=True)
47
48     959.559 MiB      0.125 MiB          df["engagement_rate"] = (df["likes"] + df["dislikes"] + df["comment_count"]) / (df["views"] + 1)
49     959.559 MiB      0.000 MiB          df["like_dislike_ratio"] = df["likes"] / (df["dislikes"] + 1)
50
51     959.809 MiB      0.250 MiB          cat_df = pd.get_dummies(df["category_id"], prefix="cat")
52     959.809 MiB      0.000 MiB          df = pd.concat([df, cat_df], axis=1)
53     959.809 MiB      0.000 MiB          df.drop(columns=["category_id"], inplace=True)
54
55     959.809 MiB      0.000 MiB          bool_cols = ["comments_disabled", "ratings_disabled", "video_error_or_removed"]
56     959.809 MiB      0.000 MiB          df[bool_cols] = df[bool_cols].astype(int)
57     959.809 MiB      0.000 MiB          df.drop(columns=bool_cols, inplace=True)
58
59     961.059 MiB      1.250 MiB          df = df.drop_duplicates().reset_index(drop=True)
60
61     961.059 MiB      0.000 MiB          numeric_attributes = [
62     961.059 MiB      0.000 MiB              "views", "publish_hour", "likes", "dislikes", "comment_count",
63     961.059 MiB      0.000 MiB              "engagement_rate", "like_dislike_ratio", "tag_count"
64     961.059 MiB      0.000 MiB          ] + [col for col in df.columns if "_emb_" in col]
65
66     961.559 MiB      0.500 MiB          df[numeric_attributes] = df[numeric_attributes].apply(lambda col: np.log1p(col))
67
68     961.559 MiB      0.000 MiB          minmax_scaler = MinMaxScaler()
69     962.059 MiB      0.500 MiB          df[numeric_attributes] = minmax_scaler.fit_transform(df[numeric_attributes])
70
71     962.059 MiB      0.000 MiB          standard_scaler = StandardScaler()
72     962.184 MiB      0.125 MiB          df[numeric_attributes] = standard_scaler.fit_transform(df[numeric_attributes])
73
74     962.184 MiB      0.000 MiB          df.drop(columns=["likes", "dislikes"], inplace=True)
75     962.184 MiB      0.000 MiB          print("Preprocessing complete. Final shape:", df.shape)
```

Before optimization, the use of *.iloc[i]* in multiple for-loops duplicated data temporarily, increasing memory overhead.

After optimization, memory-intensive manual deduplication and transformation loops were replaced with pandas-native `drop_duplicates()` and `apply(lambda col: ...)`, reducing memory peaks.

Py-spy:

```
Total Samples 6100
Gil: 1.00%, Active: 100.00%, Threads: 2
```

%Down	%Total	OwnTime	TotalTime	Function (filename)
69.00%	69.00%	39.21s	39.21s	forward (torch/nn/modules/linear.py)
23.00%	100.00%	7.09s	51.14s	forward (transformers/models/bert/modeling_bert.py)
5.00%	5.00%	2.81s	2.81s	forward (transformers/activations.py)
0.00%	0.00%	0.830s	0.830s	layer_norm (torch/nn/functional.py)
0.00%	0.00%	0.820s	0.850s	read (pandas/io/parsers/c_parser_wrapper.py)
0.00%	0.00%	0.750s	0.750s	get_data (<frozen importlib._bootstrap_external>)
2.00%	57.00%	0.570s	30.92s	apply_chunking_to_forward (transformers/pytorch_utils.py)
0.00%	0.00%	0.420s	0.420s	_compile_bytecode (<frozen importlib._bootstrap_external>)
0.00%	0.00%	0.400s	3.51s	_call_with_frames_removed (<frozen importlib._bootstrap>)
0.00%	0.00%	0.350s	0.600s	create_import_structure_from_path (transformers/utils/import_utils.py)
0.00%	0.00%	0.220s	0.220s	fetch_all (transformers/utils/import_utils.py)
0.00%	0.00%	0.200s	0.320s	_batch_encode_plus (transformers/tokenization_utils_fast.py)
0.00%	0.00%	0.160s	0.160s	embedding (torch/nn/functional.py)
0.00%	0.00%	0.130s	0.130s	_create_fn (dataclasses.py)
0.00%	0.00%	0.120s	0.120s	_expand_mask (transformers/modeling_attn_mask_utils.py)
0.00%	0.00%	0.110s	0.110s	read (ssl.py)
1.00%	100.00%	0.100s	51.23s	_call_impl (torch/nn/modules/module.py)
0.00%	0.00%	0.100s	0.100s	_path_hooks (<frozen importlib._bootstrap_external>)
0.00%	0.00%	0.090s	0.090s	__setattr__ (enum.py)
0.00%	0.00%	0.080s	0.080s	inner (tqdm/utils.py)
0.00%	0.00%	0.080s	0.080s	as_tensor (transformers/tokenization_utils_base.py)
0.00%	0.00%	0.080s	0.080s	transpose_for_scores (transformers/models/bert/modeling_bert.py)
0.00%	0.00%	0.080s	0.080s	_path_stat (<frozen importlib._bootstrap_external>)
0.00%	0.00%	0.070s	0.070s	forward (sentence_transformers/models/Pooling.py)
0.00%	0.00%	0.060s	0.070s	_has_script_object_arg (torch/_ops.py)
0.00%	0.00%	0.060s	0.060s	_joinrealpath (posixpath.py)
0.00%	0.00%	0.050s	0.050s	exists (genericpath.py)
0.00%	0.00%	0.050s	0.050s	forward (torch/nn/modules/activation.py)
0.00%	0.00%	0.040s	0.040s	decode (codecs.py)
0.00%	0.00%	0.040s	0.090s	__init__ (torch/_ops.py)
0.00%	0.00%	0.040s	0.040s	ssl_wrap_socket (urllib3/util/ssl_.py)
0.00%	0.00%	0.030s	0.030s	_fill_cache (<frozen importlib._bootstrap_external>)
0.00%	0.00%	0.030s	0.030s	docformat (cctpy/_lib/doccer.py)
0.00%	0.00%	0.030s	0.030s	<listcomp> (<frozen importlib._bootstrap_external>)
0.00%	0.00%	0.030s	0.030s	dedent (textwrap.py)
0.00%	0.00%	0.030s	0.030s	impl (torch/library.py)
0.00%	0.00%	0.030s	0.030s	raw_decode (json/decoder.py)
0.00%	0.00%	0.030s	0.030s	_prepare_class_assumptions (sympy/core/assumptions.py)
0.00%	0.00%	0.030s	0.040s	__init__ (inspect.py)
0.00%	0.00%	0.030s	0.030s	open (pathlib.py)
0.00%	0.00%	0.030s	0.030s	isdir (genericpath.py)
0.00%	0.00%	0.030s	0.030s	_get_packet (torch/_ops.py)
0.00%	0.00%	0.030s	0.280s	_process_class (dataclasses.py)
0.00%	0.00%	0.020s	3.87s	_find_and_load_unlocked (<frozen importlib._bootstrap>)
0.00%	0.00%	0.020s	0.040s	<listcomp> (sentence_transformers/SentenceTransformer.py)
0.00%	0.00%	0.020s	0.030s	set_truncation_and_padding (transformers/tokenization_utils_fast.py)
0.00%	0.00%	0.020s	0.020s	vstack (numpy/_core/shape_base.py)
0.00%	0.00%	0.020s	0.030s	__init__ (transformers/utils/import_utils.py)
0.00%	0.00%	0.020s	0.020s	dropout (torch/nn/functional.py)
0.00%	0.00%	0.020s	0.020s	__init__ (pandas/io/parsers/c_parser_wrapper.py)

Press **Control-C** to quit, or **?** for help.

sentence encoding is expensive and stays the core bottleneck.

Part 3: compare three different methods for applying a condition to a column in a DataFrame

suppose we have a column called "views" and we want to create a new column "popularity" that contains "popular" if *views* > 100,000 and "not popular" otherwise.

1. Using pandas apply:

```
df["popularity"] = df["views"].apply(lambda x: "popular" if x > 100000 else "not popular")
```

This is the slowest method because it loops over rows in Python. However, it is the most readable one.

2. Using pandas map

```
df["popularity"] = (df["views"] > 100000).map({True: "popular", False: "not popular"})
```

This is faster than apply but less readable.

3. Using numpy where

```
df["popularity"] = np.where(df["views"] > 100000, "popular", "not popular")
```

This is the fastest and most memory-efficient due to vectorized operations but the least readable.

Execution time, memory usage, CPU usage, and readability are the key metrics to compare.

Overall, if performance matters (especially on large datasets), prefer `np.where`. If clarity is more important for the reader or for debugging, `apply` may be a better fit.

Part 4: Questions

۱. پروفایلینگ در محیط Production باید با دقت بیشتری انجام شود تا کمترین تأثیر را روی عملکرد سیستم داشته باشد و اطلاعات حساس در معرض خطر قرار نگیرند. معمولاً از ابزارهای سبک مثل py-spy یا cProfile در حالت نمونه‌برداری استفاده می‌شود چون می‌توانند بدون توقف برنامه، آن را بررسی کنند. در این محیط، باید به تأثیر عملکردی، امنیت و دقت داده‌ها در شرایط واقعی توجه ویژه داشت.

۲. در برنامه‌های multithreaded، پروفایلینگ پیچیده‌تر است چون زمان اجرا و مصرف CPU بین چند نخ پخش می‌شود. ابزارهایی مثل py-spy یا perf برای این شرایط مناسب‌تر هستند چون می‌توانند عملکرد هر نخ را جداگانه بررسی کنند. برخلاف برنامه‌های تک‌ریسمانی، در اینجا باید به مسائل مربوط به رقابت بین نخ‌ها، جابه‌جایی بین نخ‌ها و محدودیت‌های GIL در پایتون توجه کرد تا تحلیل دقیقی به دست آورد.